



Adrien Bertaud <bertaud.adrien@gmail.com>

(aucun objet)

Viano Luca <luca.viano@epfl.ch>

21 juillet 2020 à 13:13

À : Simsek Berfin <berfin.simsek@epfl.ch>, Adrien Bertaud <bertaud.adrien@gmail.com>

Cc : Müller Nathan Samuel <nathan.muller@epfl.ch>, Kuras Ihor <ihor.kuras@epfl.ch>

Hello Adrien and Nathan,

Congrats also from my side!

As far as your error in exercise 1 are concerned, I penalized you for some answers to the theoretical questions that were not fully correct. In particular:

In question 2 I was expecting an answer pointing to the fact that the formulation of the REINFORCE algorithm use the gradient computed for one sample as an unbiased estimator for the expected gradient. If you use also past data, then you get a biased estimator. Your answer is partially correct but a bit vague about this bias problem.

Finally, in question 4 you forgot to mention that you should modify the loss from binary cross entropy to categorical cross entropy.

Best regards,

Luca

Da: simsek <berfin.simsek@epfl.ch>**Inviato:** martedì 21 luglio 2020 10:50:48**A:** Adrien Bertaud**Cc:** Müller Nathan Samuel; Viano Luca; Kuras Ihor**Oggetto:** Re:

[Texte des messages précédents masqué]